

Assignment 2: Q2 (System Call)

SHREY MARWAHA – 2019334

Adding A System Call to The Linux Kernel (5.9.1) In Ubuntu (20.04.1 LTS)

Part 1: Preparation

The linux 5.9.1 source code is downloaded and it's extracted using the following.

- 'wget -P ~/ <https://cdn.kernel.org/pub/linux/kernel/v5.x/linux5.9.1.tar.xz>'
- A new directory `sh_task_info` is made which contains 2 files – `sh_task_info.c` and `Makefile`.
- `sh_task_info.c` – functions take `int pid` and `char* filename` as parameters and uses `write` to write it to a file. In `SYSCALL_DEFINE2(sh_task_info, int, pid, char*, filename)` iterates over all the process by using `'for_each_process()'` and compares the pid of each process with the given pid.
- `Makefile` contains the following code: `'obj-y := sh_task_info.o'`
- In the root of folder linux 5.9.1 we edit `Makefile` and append `'sh_task_info/'` to the `core-y:`
- In the file `nano include/linux/syscalls.h` we Add a corresponding function prototype for your system call to the header file of system calls just above `#endif` in the bottom.
- Add system call to the kernel's system call table by `'nano arch/x86/entry/syscalls/syscall_64.tbl'`

Part 2: Compilation and Installation

- First a `.config` file is created with command `'make menuconfig'`.
- It was compiled with the three parallel commands by `'sudo make -j4 && sudo make modules_install -j4 && sudo make install -j4'`.
- Reboot the machine after this step.

Part 3: Testing

- The file `'report.c'` is the sample test code, compares the process IDs by using `syscall`. If the input process ID and the updated pid match, and there is no error in the given pid and the filename `sh_task_info()` is executed correctly.
- It should be compiled with `'gcc report.c -o report'`

```
shrey@shrey-VirtualBox:~$ gcc report.c -o report
```

- The executable is run with './report' and provide the **PID** and **FILENAME**

```
shrey@shrey-VirtualBox:~$ ./report

Enter the pid: 10
Enter the file name: /home/shrey/Desktop/output.txt

Congratulations! Your system call is functional.
Run the command dmesg in the terminal and find out!

shrey@shrey-VirtualBox:~$ dmesg
```

- If the execution is successful, then we can check the kernel log via 'dmesg'. The attributes corresponding to pid will be printed on the console.

- The following attributes are printed in **kernel log** and printed in **the file** provided:-

1. PROCESS
2. PID NUMBER
3. PROCESS STATE
4. PRIORITY
5. RT_PRIORITY
6. STATIC PRIORITY
7. NORMAL PRIORITY

```
[ 5861.705897]
                PROCESS: ksoftirqd/0
                PID_NUMBER: 10
                PROCESS STATE: 1
                PRIORITY: 120
                RT_PRIORITY: 0
                STATIC PRIORITY: 120
                NORMAL PRIORITY: 120
shrey@shrey-VirtualBox:~$
```

- Alternatively, we can check the log with '**cat <FILENAME>.txt**'.

```
1
2 Process: ksoftirqd/0
3 PID_NUMBER: 10
4 PROCESS STATE: 1
5 PRIORITY: 120
6 RT_PRIORITY: 0
7 STATIC PRIORITY: 120
8 NORMAL PRIORIY: 120
```

Part 4: Errors Handled

- i. If the user enters invalid pid data type, ie. anyother data type other than int such as float or char etc

```
shrey@shrey-VirtualBox:~$ ./report
Enter the pid: 5
Enter the file name: shr.txt
Sorry,system call appears to have failed.: No such process
shrey@shrey-VirtualBox:~$
```

- ii. If the user enters pid ≤ 0 or greater than 32768, the function returns the errno 22 EINVAL invalid argument.

```
shrey@shrey-VirtualBox:~$ ./report
Enter the pid: 0
Enter the file name: output.txt
Sorry,system call appears to have failed.: Invalid argument
shrey@shrey-VirtualBox:~$
```

- iii. If the file to be written does not exist then then, it will automatically creates a new file with the provided name in the /home directory.
- iv. User has the choice to enter directly the directory rather than entering the whole directory in the **filename**
Eg. instead of providing filename as *'/home/shrey/Desktop/output.txt'* user has the choice to enter as *'Desktop/output.txt'*

```
shrey@shrey-VirtualBox:~$ ./report
Enter the pid: 1
Enter the file name: Desktop/shr.txt

Congratulations! Your system call is functional.
Run the command dmesg in the terminal and find out!
shrey@shrey-VirtualBox:~$
```

- v. If user enters wrong directory address or any directory that doesn't even exist then the program will show **'Killed'**

```
shrey@shrey-VirtualBox:~$ ./report
Enter the pid: 10
Enter the file name: /home/Desktop/output.txt
Killed
shrey@shrey-VirtualBox:~$
```

- vi. If user doesn't provide any directory for file then the program will by default take the directory as **home directory**.